

# Operating Systems

## CMPSCI 377

### Spring 2020

# Clicker Question

What is the best way to describe how user programs run?

- (A) directly on the CPU
- (B) on top of the Operating System, which runs on top of the CPU
- (C) inside of the OS
- (D) on the GPU

Answer on Next Slide

# Today's Class

- Limited Direct Execution

## THE CRUX:

### HOW TO EFFICIENTLY VIRTUALIZE THE CPU WITH CONTROL

The OS must virtualize the CPU in an efficient manner while retaining control over the system. To do so, both hardware and operating-system support will be required. The OS will often use a judicious bit of hardware support in order to accomplish its work effectively.

# Basic Technique: Limited Direct Execution

To make a program run as fast as one might expect, not surprisingly OS developers came up with a technique, which we call:

**limited direct execution.**

The direct execution part of the idea is simple: just run the program directly on the CPU.

OS	Program
Create entry for process list	
Allocate memory for program	
Load program into memory	
Set up stack with argc/argv	
Clear registers	
Execute <b>call</b> main()	Run main()
	Execute <b>return</b> from main
Free memory of process	
Remove from process list	

# Basic Technique: Limited Direct Execution

To make a program run as fast as one might expect, not surprisingly OS developers came up with a technique, which we call:

**limited direct execution.**

The direct execution part of the idea is simple: just run the program directly on the CPU.

OS

Create entry for process list  
Allocate memory for program  
Load program into memory  
Set up stack with argc/argv  
Clear registers  
Execute **call** main()

Free memory of process  
Remove from process list

Program

Run main()  
Execute **return** from main

# Basic Technique: Limited Direct Execution

To make a program run as fast as one might expect, not surprisingly OS developers came up with a technique, which we call:

**limited direct execution.**

The direct execution part of the idea is simple: just run the program directly on the CPU.

OS	Program
Create entry for process list Allocate memory for program Load program into memory Set up stack with argc/argv Clear registers Execute <b>call</b> main()	Run main() Execute <b>return</b> from main
Free memory of process Remove from process list	



# Basic Technique: Limited Direct Execution

To make a program run as fast as one might expect, not surprisingly OS developers came up with a technique, which we call:

**limited direct execution.**

The direct execution part of the idea is simple: just run the program directly on the CPU.

## OS

- Create entry for process list
- Allocate memory for program
- Load program into memory
- Set up stack with argc/argv
- Clear registers
- Execute **call** main()

- Free memory of process
- Remove from process list

## Program

- Run main()
- Execute **return** from main

# Basic Technique: Limited Direct Execution

Pretty simple...

But, there are problems...

OS

Program

---

Create entry for process list  
Allocate memory for program  
Load program into memory  
Set up stack with argc/argv  
Clear registers  
Execute **call** main()

Run main()  
Execute **return** from main

Free memory of process  
Remove from process list

# Basic Technique: Limited Direct Execution

## Problem #1

### Restricted Operations

If we just run a program, how can the OS make sure the program doesn't do anything that we don't want it to do, while still running it efficiently?

#### OS

---

- Create entry for process list
- Allocate memory for program
- Load program into memory
- Set up stack with argc/argv
- Clear registers
- Execute **call** main()

- Free memory of process
- Remove from process list

#### Program

---

- Run main()
- Execute **return** from main

# Basic Technique: Limited Direct Execution

## Problem #2

### Context Switching

When we are running a process, how does the operating system stop it from running and switch to another process, thus implementing the time sharing we require to virtualize the CPU?

OS	Program
Create entry for process list	
Allocate memory for program	
Load program into memory	
Set up stack with argc/argv	
Clear registers	
Execute <b>call</b> main()	Run main()
	Execute <b>return</b> from main
Free memory of process	
Remove from process list	

# Problem #1: Restricted Operations

## THE CRUX: HOW TO PERFORM RESTRICTED OPERATIONS

A process must be able to perform I/O and some other restricted operations, but without giving the process complete control over the system. How can the OS and hardware work together to do so?

# Here is a bad solution

We could let a process do whatever it wants.

# Clicker Question

If there are no restrictions on the user process, what can it do?

- (A) corrupt the memory of the kernel
- (B) halt the machine
- (C) run forever and starve other processes
- (D) all of the above

Answer on Next Slide



# Here is a bad solution

We could let a process do whatever it wants.

But, this is likely to be a very bad idea.

It would prevent the construction of many kinds of systems that are desirable.



For example, we can't simply let any user process issue I/Os to disk.

If we did, a process could simply read or write an entire disk and thus all protections would be lost.


# Introducing: User Mode

Thus, the approach we take is to introduce a new processor mode known as user mode.

Code that runs in user mode is restricted in what it can do.

For example, no I/O requests!

If you try, then




# Introducing: Kernel Mode

And the other processor mode  
is known as kernel mode.

Code that runs in kernel mode can do  
whatever it wants.

For example, I/O requests!

If you try, then



# Clicker Question

True or False: A user program can do something to change the machine to kernel mode.

(A) True

(B) False

Answer on Next Slide

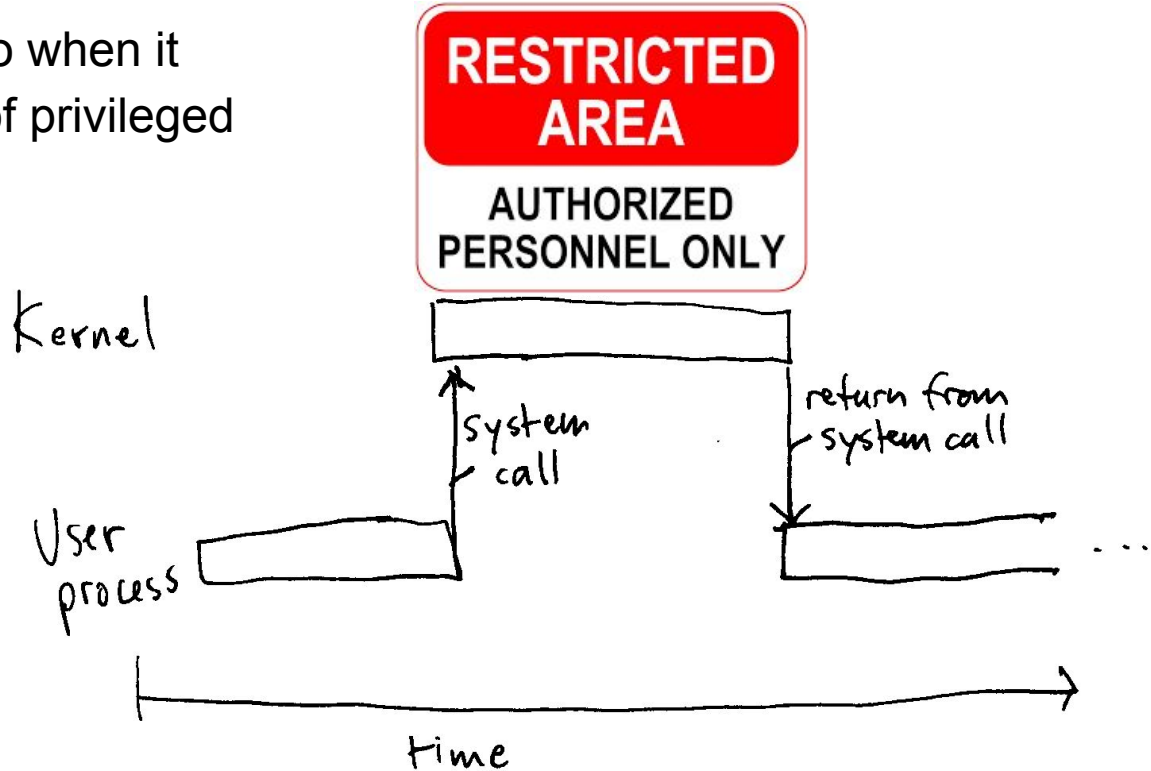
# Question?

What should a user process do when it wishes to perform some kind of privileged operation?

# Question?

What should a user process do when it wishes to perform some kind of privileged operation?

To enable this, virtually all modern hardware provides the ability for user programs to perform a system call.



# System Calls

System calls expose key pieces of functionality to user programs

- Accessing the file system
- Creating and destroying processes
- Communicating with other processes
- Allocating more memory

Most operating systems provide a few hundred system calls

- POSIX
- Early Unix systems had a more concise set of 20



# Question?

If the kernel allows privileged functionality, how does a system call get us into the kernel?

How are system calls implemented?

We will describe this process in a “general” way.

There are a lot of specifics and differences between x86 and x86-64

Read: <https://blog.packagecloud.io/eng/2016/04/05/the-definitive-guide-to-linux-system-calls/>

# Executing a System Call

To execute a system call:

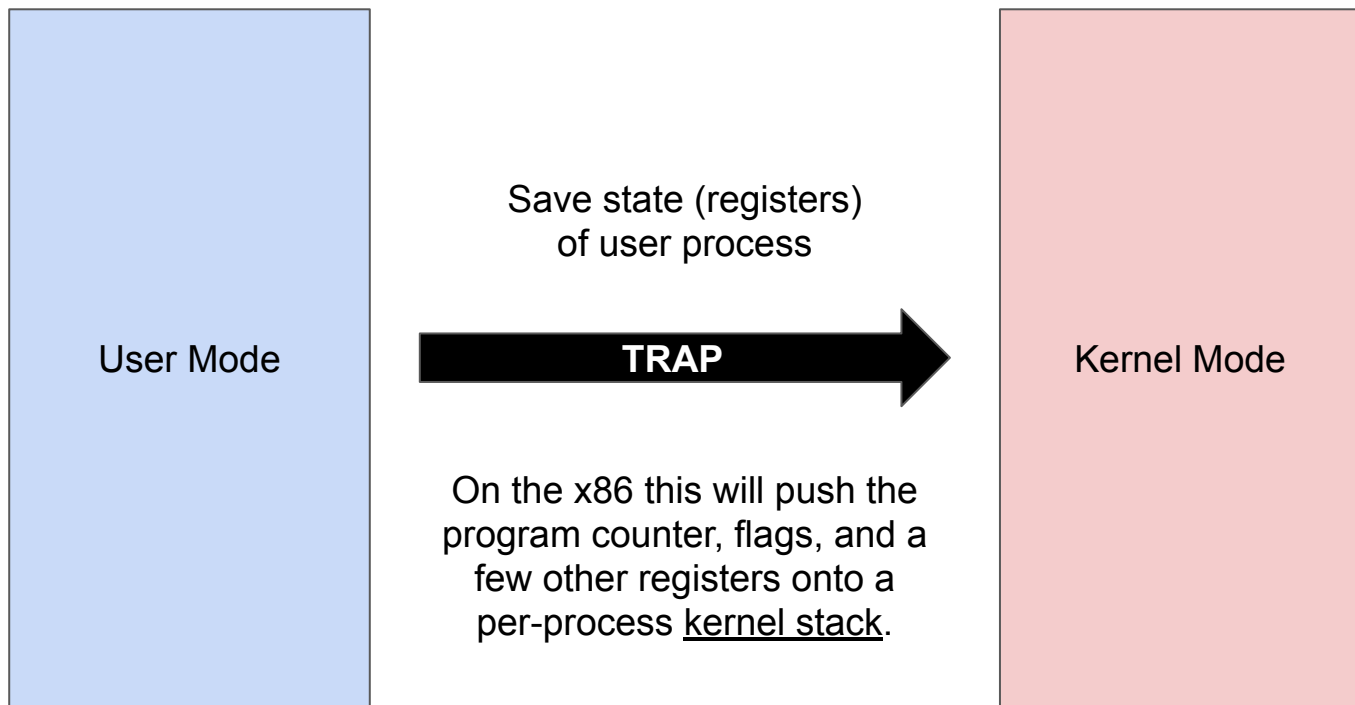
- A program executes a special trap instruction
  - Jumps into the kernel and raises the privilege level at the same time
  - Privileged operations **can be** performed

# Executing a System Call

To execute a system call:

- A program executes a special trap instruction
  - Jumps into the kernel and raises the privilege level at the same time
  - Privileged operations **can be** performed
- When finished, the OS calls a special return-from-trap instruction
  - Returns to the calling program and lowers the privilege level at the same time
  - Privileged operations **can not be** performed

# Preparing for a System Call: Saving State

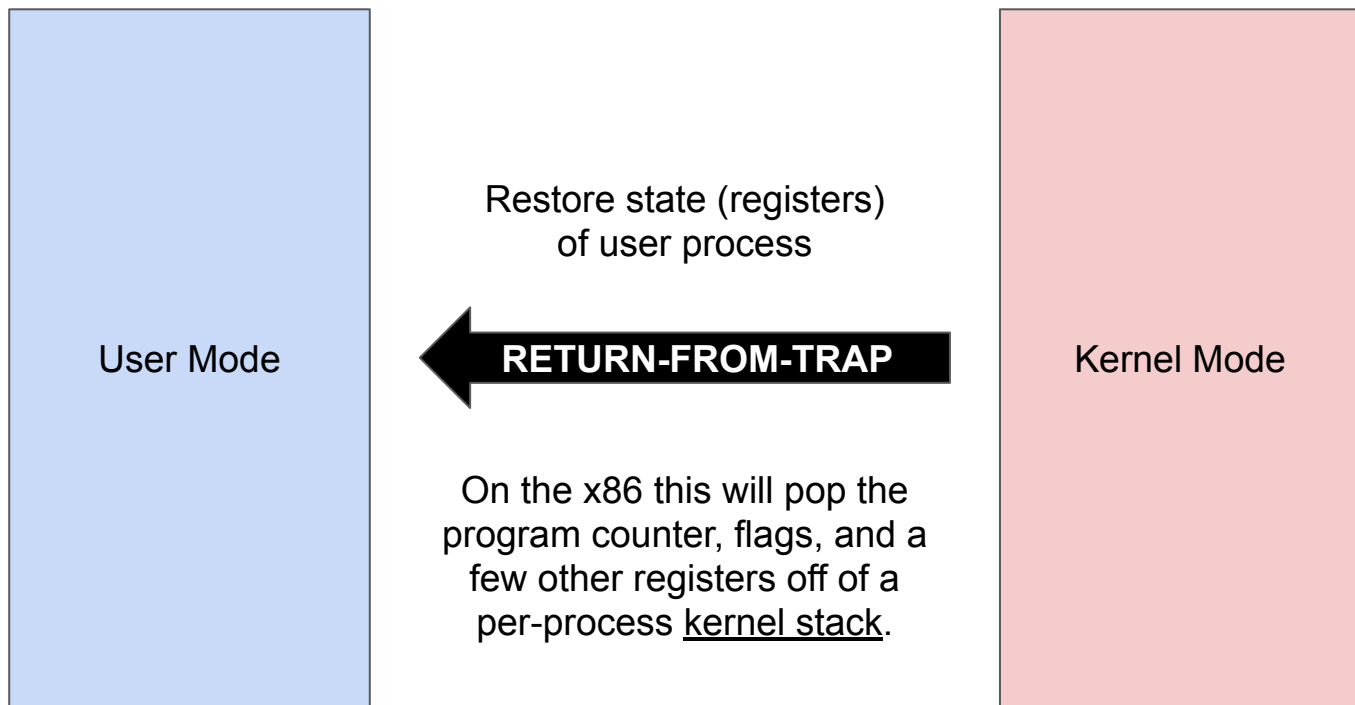


# Preparing for a System Call: Saving State

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                              // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If non-zero, sleeping on chan
    int killed;               // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;         // Current directory
    struct context *ctx;       // Switch here to run process
    struct trapframe *tf;      // Trap frame for the
                              // current interrupt
};
```

Remember this from last time?

# Completing a System Call: Restoring State



How does the trap know  
which code to run in the OS?



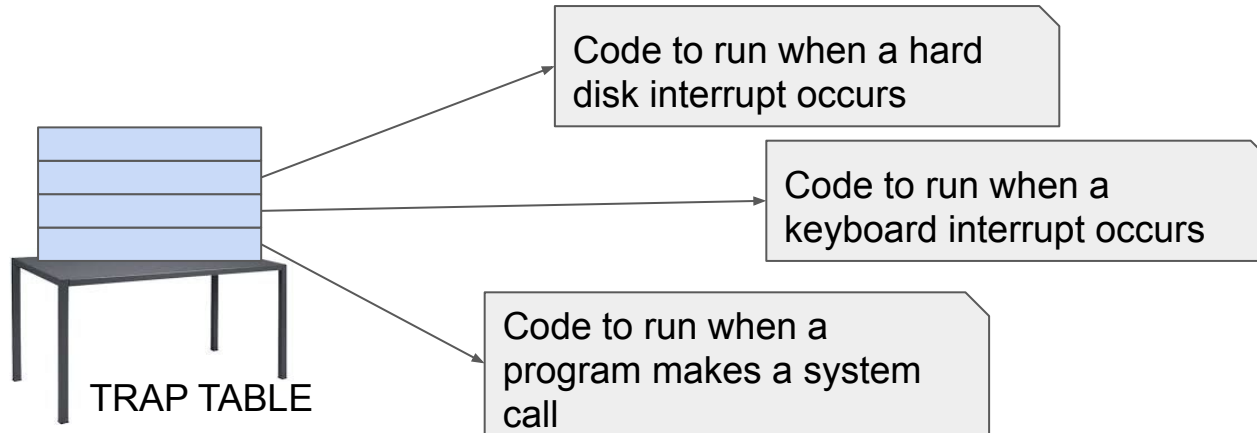
Can we simply use an address?

# At boot time...



When the machine boots up, it does so in privileged (kernel) mode, and thus is free to configure machine hardware as need be.

One of the first things the OS thus does is to tell the hardware what code to run when certain exceptional events occur (special instruction).






# At boot time...



When the machine boots up, it does so in privileged (kernel) mode, and thus is free to configure machine hardware as need be.

One of the first things the OS thus does is to tell the hardware what code to run when certain exceptional events occur (special instruction).



Once the hardware is informed, it remembers the location of these handlers until the machine is next rebooted, and thus the hardware knows what to do (i.e., what code to jump to) when system calls and other exceptional events take place.

“All problems in computer science can be solved by  
another level of indirection”

- David Wheeler

# System Call Numbers

To specify the exact system call, a system call number is usually assigned to each system call.

“All problems in computer science can be solved by  
another level of indirection”

- David Wheeler

# System Call Numbers

To specify the exact system call, a system call number is usually assigned to each system call.

The user code is thus responsible for placing the desired system call number in a register or at a specified location on the stack.

“All problems in computer science can be solved by  
another level of indirection”

- David Wheeler

# System Call Numbers

To specify the exact system call, a system call number is usually assigned to each system call.

The user code is thus responsible for placing the desired system call number in a register or at a specified location on the stack.

The OS, when handling the system call inside the trap handler, examines this number, ensures it is valid, and, if it is, executes the corresponding code.

“All problems in computer science can be solved by  
another level of indirection”

- David Wheeler

# System Call Numbers

To specify the exact system call, a system call number is usually assigned to each system call.

The user code is thus responsible for placing the desired system call number in a register or at a specified location on the stack.

The OS, when handling the system call inside the trap handler, examines this number, ensures it is valid, and, if it is, executes the corresponding code.

This level of **indirection** serves as a form of **protection**; user code cannot specify an exact address to jump to, but rather must request a particular service via number.

**OS @ boot  
(kernel mode)**

**Hardware**

---

**initialize trap table**

remember address of...  
syscall handler

**OS @ run**  
**(kernel mode)**

**Hardware**

**Program**  
**(user mode)**

---

Limited Direct Execution Protocol

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC

Limited Direct Execution Protocol



**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to main

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
<b>return-from-trap</b>		
	restore regs from kernel stack	
	move to user mode	
	jump to main	
		Run main()
		...
		Call system call
		<b>trap</b> into OS

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
<b>return-from-trap</b>		
	restore regs from kernel stack	
	move to user mode	
	jump to main	
		Run main()
		...
		Call system call
		<b>trap</b> into OS
	save regs to kernel stack	
	move to kernel mode	
	jump to trap handler	

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
<b>return-from-trap</b>		
	restore regs from kernel stack	
	move to user mode	
	jump to main	
		Run main()
		...
		Call system call
		<b>trap</b> into OS
	save regs to kernel stack	
	move to kernel mode	
	jump to trap handler	
Handle trap		
Do work of syscall		
<b>return-from-trap</b>		

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
<b>return-from-trap</b>		
	restore regs from kernel stack	
	move to user mode	
	jump to main	
		Run main()
		...
		Call system call
		<b>trap</b> into OS
	save regs to kernel stack	
	move to kernel mode	
	jump to trap handler	
Handle trap		
Do work of syscall		
<b>return-from-trap</b>		
	restore regs from kernel stack	
	move to user mode	
	jump to PC after trap	

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC <b>return-from-trap</b>		
	restore regs from kernel stack move to user mode jump to main	Run main() ... Call system call <b>trap</b> into OS
	save regs to kernel stack move to kernel mode jump to trap handler	
Handle trap Do work of syscall <b>return-from-trap</b>	restore regs from kernel stack move to user mode jump to PC after trap	... return from main <b>trap</b> (via <code>exit()</code> )

Limited Direct Execution Protocol

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
<b>return-from-trap</b>	restore regs from kernel stack move to user mode jump to main	Run main() ... Call system call <b>trap</b> into OS
	save regs to kernel stack move to kernel mode jump to trap handler	
Handle trap		
Do work of syscall		
<b>return-from-trap</b>	restore regs from kernel stack move to user mode jump to PC after trap	... return from main <b>trap</b> (via <code>exit()</code> )
Free memory of process		
Remove from process list		

Limited Direct Execution Protocol

# Kernel Stack vs User Space Stack?

<https://stackoverflow.com/questions/12911841/kernel-stack-and-user-space-stack>



142



1. *What's the difference between kernel stack and user stack ?*

In short, nothing - apart from using a different location in memory (and hence a different value for the stackpointer register), and usually different memory access protections. I.e. when executing in user mode, kernel memory (part of which is the kernel stack) will not be accessible even if mapped. Vice versa, without explicitly being requested by the kernel code (in Linux, through functions like `copy_from_user()`), user memory (including the user stack) is not usually directly accessible.

2. *Why is [ a separate ] kernel stack used ?*

Separation of privileges and security. For one, userspace programs can make their stack(pointer) anything they want, and there is usually no architectural requirement to even have a valid one. The kernel therefore cannot *trust* the userspace stackpointer to be valid nor usable, and therefore will require one set under its own control. Different CPU architectures implement this in different ways; x86 CPUs automatically switch stackpointers when privilege mode switches occur, and the values to be used for different privilege levels are configurable - by privileged code (i.e. only the kernel).



# Problem #2: Switching Between Processes

THE CRUX: HOW TO REGAIN CONTROL OF THE CPU

How can the operating system **regain control** of the CPU so that it can switch between processes?

# A Cooperative Approach: Wait for System Calls

in a cooperative scheduling system, the OS regains control of the CPU by waiting for a system call or an illegal operation of some kind to take place.



# A Non-Cooperative Approach: OS Takes Control

THE CRUX: HOW TO GAIN CONTROL WITHOUT COOPERATION

How can the OS gain control of the CPU even if processes are not being cooperative? What can the OS do to ensure a rogue process does not take over the machine?



# A Non-Cooperative Approach: OS Takes Control

## THE CRUX: HOW TO GAIN CONTROL WITHOUT COOPERATION

How can the OS gain control of the CPU even if processes are not being cooperative? What can the OS do to ensure a rogue process does not take over the machine?

## TIP: USE THE TIMER INTERRUPT TO REGAIN CONTROL

The addition of a **timer interrupt** gives the OS the ability to run again on a CPU even if processes act in a non-cooperative fashion. Thus, this hardware feature is essential in helping the OS maintain control of the machine.

# Saving and Restoring Context

Now that the OS has regained control, whether cooperatively via a system call, or more forcefully via a timer interrupt, a decision has to be made:

- whether to continue running the currently-running process
- or switch to a different one.

# Saving and Restoring Context

Now that the OS has regained control, whether cooperatively via a system call, or more forcefully via a timer interrupt, a decision has to be made:

- whether to continue running the currently-running process
- or switch to a different one.

This decision is made by a part of the operating system known as the scheduler  
we will discuss scheduling policies in detail next week.



Scheduler

If the decision is made to **switch**, the OS then executes a low-level piece of code which we refer to as a context switch.

CONTEXT  
MATTERS



## OS @ boot (kernel mode)

## Hardware

---

**initialize trap table**

remember addresses of...  
syscall handler  
timer handler

**start interrupt timer**

start timer  
interrupt CPU in X ms



**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Process A

...

Limited Direct Execution Protocol (timer interrupt)

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Process A

...

**timer interrupt**

save regs(A) to k-stack(A)

move to kernel mode

jump to trap handler

Limited Direct Execution Protocol (timer interrupt)

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Process A

...

**timer interrupt**

save regs(A) to k-stack(A)

move to kernel mode

jump to trap handler

Handle the trap

Call `switch()` routine

save regs(A) to proc-struct(A)

restore regs(B) from proc-struct(B)

switch to k-stack(B)

**return-from-trap (into B)**

Limited Direct Execution Protocol (timer interrupt)

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Process A

...

**timer interrupt**

save regs(A) to k-stack(A)

move to kernel mode

jump to trap handler

Handle the trap

Call `switch()` routine

save regs(A) to proc-struct(A)

restore regs(B) from proc-struct(B)

switch to k-stack(B)

**return-from-trap (into B)**

restore regs(B) from k-stack(B)

move to user mode

jump to B's PC

Limited Direct Execution Protocol (timer interrupt)

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Process A

...

**timer interrupt**

save regs(A) to k-stack(A)

move to kernel mode

jump to trap handler

Handle the trap

Call `switch()` routine

save regs(A) to proc-struct(A)

restore regs(B) from proc-struct(B)

switch to k-stack(B)

**return-from-trap (into B)**

restore regs(B) from k-stack(B)

move to user mode

jump to B's PC

Process B


...

Limited Direct Execution Protocol (timer interrupt)

Hmm... what happens when, during a system call, a timer interrupt occurs?" or "What happens when you're handling one interrupt and another one happens?



Hmm... what happens when, during a system call, a timer interrupt occurs?” or “What happens when you’re handling one interrupt and another one happens?



Hmm... what happens when, during a system call, a timer interrupt occurs?” or “What happens when you’re handling one interrupt and another one happens?





# Clicker Question

When a timer interrupt occurs causing a context switch from process A to B, what is missing here?

Use Process A running -> Timer interrupt occurs -> CPU save user registers in kernel stack -> Kernel Code -> Kernel saves kernel registers -> Kernel switches to another kernel thread -> <MISSING> -> Process B runs

(A) CPU saves registers to kernel stack

(B) Kernel invokes return-from-trap

(C) Process B calls a trap instruction

(D) The CPU halts

Answer on Next Slide

# Next Time

CPU Scheduling